

In []:

In []:

In []:

```
In [1]: from tensorflow import lite
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
import numpy as np
import pandas as pd
import random, os
import shutil
import matplotlib.pyplot as plt
from matplotlib.image import imread
# from keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.metrics import categorical_accuracy
from sklearn.model_selection import train_test_split
```

```
In [4]: # Add an additional column, mapping to the type
df = pd.read_csv(r'train.csv')

diagnosis_dict_binary = {
    0: 'No_DR',
    1: 'DR',
    2: 'DR',
    3: 'DR',
    4: 'DR'
}

diagnosis_dict = {
    0: 'No_DR',
    1: 'Mild',
    2: 'Moderate',
    3: 'Severe',
    4: 'Proliferate_DR',
}


df['binary_type'] = df['diagnosis'].map(diagnosis_dict_binary.get)
df['type'] = df['diagnosis'].map(diagnosis_dict.get)
df.head()
```

Out[4]:

	id_code	diagnosis	binary_type	type
0	000c1434d8d7	2	DR	Moderate
1	001639a390f0	4	DR	Proliferate_DR
2	0024cdab0c1e	1	DR	Mild
3	002c21358ce6	0	No_DR	No_DR
4	005b95c28852	0	No_DR	No_DR


```
In [5]: df['type'].value_counts().plot(kind='barh')
```

Out[5]: <Axes: ylabel='type'>



```
In [6]: df['binary_type'].value_counts().plot(kind='barh')
```

Out[6]: <Axes: ylabel='binary_type'>



```
In [7]: # Split into stratified train, val, and test sets
train_intermediate, val = train_test_split(df, test_size = 0.15, stratify = df['type'],
train, test = train_test_split(train_intermediate, test_size = 0.15 / (1 - 0.15),

print(train['type'].value_counts(), '\n')
```

```
print(test['type'].value_counts(), '\n')
print(val['type'].value_counts(), '\n')
```

```
type
No_DR          1263
Moderate        699
Mild            258
Proliferate_DR  207
Severe          135
Name: count, dtype: int64
```

```
type
No_DR          271
Moderate        150
Mild            56
Proliferate_DR  44
Severe          29
Name: count, dtype: int64
```

```
type
No_DR          271
Moderate        150
Mild            56
Proliferate_DR  44
Severe          29
Name: count, dtype: int64
```

In [8]: *# Create working directories for train/val/test*

```
base_dir = ''

train_dir = os.path.join(base_dir, 'train')
val_dir = os.path.join(base_dir, 'val')
test_dir = os.path.join(base_dir, 'test')

if os.path.exists(base_dir):
    shutil.rmtree(base_dir)

if os.path.exists(train_dir):
    shutil.rmtree(train_dir)
os.makedirs(train_dir)

if os.path.exists(val_dir):
    shutil.rmtree(val_dir)
os.makedirs(val_dir)

if os.path.exists(test_dir):
    shutil.rmtree(test_dir)
os.makedirs(test_dir)
```

In [12]: *# Copy images to respective working directory*

```
src_dir = r'gaussian_filtered_images\gaussian_filtered_images'
for index, row in train.iterrows():
    diagnosis = row['type']
    binary_diagnosis = row['binary_type']
    id_code = row['id_code'] + ".png"
    srcfile = os.path.join(src_dir, diagnosis, id_code)
    dstfile = os.path.join(train_dir, binary_diagnosis)
    os.makedirs(dstfile, exist_ok = True)
    shutil.copy(srcfile, dstfile)
```

```

for index, row in val.iterrows():
    diagnosis = row['type']
    binary_diagnosis = row['binary_type']
    id_code = row['id_code'] + ".png"
    srcfile = os.path.join(src_dir, diagnosis, id_code)
    dstfile = os.path.join(val_dir, binary_diagnosis)
    os.makedirs(dstfile, exist_ok = True)
    shutil.copy(srcfile, dstfile)

for index, row in test.iterrows():
    diagnosis = row['type']
    binary_diagnosis = row['binary_type']
    id_code = row['id_code'] + ".png"
    srcfile = os.path.join(src_dir, diagnosis, id_code)
    dstfile = os.path.join(test_dir, binary_diagnosis)
    os.makedirs(dstfile, exist_ok = True)
    shutil.copy(srcfile, dstfile)

```

```

In [13]: # Setting up ImageDataGenerator for train/val/test
from tensorflow.keras.preprocessing.image import ImageDataGenerator

train_path = 'train'
val_path = 'val'
test_path = 'test'

train_batches = ImageDataGenerator(rescale = 1./255).flow_from_directory(train_path)
val_batches = ImageDataGenerator(rescale = 1./255).flow_from_directory(val_path)
test_batches = ImageDataGenerator(rescale = 1./255).flow_from_directory(test_path)

```

Found 2562 images belonging to 2 classes.

Found 550 images belonging to 2 classes.

Found 550 images belonging to 2 classes.

```

In [14]: # Building the model

model = tf.keras.Sequential([
    # Layer 1: Convolutional Layer
    layers.Conv2D(8, (3,3), padding="valid", input_shape=(224,224,3), activation = 'relu'),
    layers.MaxPooling2D(pool_size=(2,2)),
    layers.BatchNormalization(),

    # Layer 2: Convolutional Layer
    layers.Conv2D(16, (3,3), padding="valid", activation = 'relu'),
    layers.MaxPooling2D(pool_size=(2,2)),
    layers.BatchNormalization(),

    # Layer 3: Convolutional Layer
    layers.Conv2D(32, (4,4), padding="valid", activation = 'relu'),
    layers.MaxPooling2D(pool_size=(2,2)),
    layers.BatchNormalization(),

    # Flatten the data
    layers.Flatten(),

    # Fully Connected Layer 1 (Dense Layer)
    layers.Dense(32, activation = 'relu'),
    layers.Dropout(0.15),

    # Output Layer
    layers.Dense(2, activation = 'softmax')
])

model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate = 1e-5),

```


```
loss=tf.keras.losses.BinaryCrossentropy(),  
metrics=['acc'])  
  
history = model.fit(train_batches,  
                    epochs=30,  
                    validation_data=val_batches)
```


c:\Users\Gakps\anaconda3\envs\final\Lib\site-packages\keras\src\layers\convolutional\base_conv.py:107: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.


```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```


c:\Users\Gakps\anaconda3\envs\final\Lib\site-packages\keras\src\trainers\data_adapters\py_dataset_adapter.py:121: UserWarning: Your `PyDataset` class should call `super().__init__(**kwargs)` in its constructor. `**kwargs` can include `workers`, `use_multiprocessing`, `max_queue_size`. Do not pass these arguments to `fit()`, as they will be ignored.


```
self._warn_if_super_not_called()
```


Epoch 1/30
81/81  **22s** 213ms/step - acc: 0.6760 - loss: 0.6624 - val_acc: 0.5073 - val_loss: 0.7078


Epoch 2/30
81/81  **16s** 201ms/step - acc: 0.8887 - loss: 0.2837 - val_acc: 0.5073 - val_loss: 0.6908


Epoch 3/30
81/81  **16s** 198ms/step - acc: 0.9176 - loss: 0.2487 - val_acc: 0.5800 - val_loss: 0.6391


Epoch 4/30
81/81  **17s** 203ms/step - acc: 0.9211 - loss: 0.2313 - val_acc: 0.6709 - val_loss: 0.5522


Epoch 5/30
81/81  **17s** 205ms/step - acc: 0.9264 - loss: 0.2149 - val_acc: 0.9109 - val_loss: 0.4127


Epoch 6/30
81/81  **17s** 203ms/step - acc: 0.9243 - loss: 0.2265 - val_acc: 0.9273 - val_loss: 0.3060


Epoch 7/30
81/81  **17s** 205ms/step - acc: 0.9377 - loss: 0.1973 - val_acc: 0.9327 - val_loss: 0.2471


Epoch 8/30
81/81  **17s** 206ms/step - acc: 0.9393 - loss: 0.2032 - val_acc: 0.9291 - val_loss: 0.2146


Epoch 9/30
81/81  **17s** 203ms/step - acc: 0.9392 - loss: 0.1934 - val_acc: 0.9436 - val_loss: 0.1846


Epoch 10/30
81/81  **17s** 203ms/step - acc: 0.9430 - loss: 0.1751 - val_acc: 0.9436 - val_loss: 0.1824


Epoch 11/30
81/81  **17s** 209ms/step - acc: 0.9509 - loss: 0.1598 - val_acc: 0.9455 - val_loss: 0.1726


Epoch 12/30
81/81  **16s** 199ms/step - acc: 0.9443 - loss: 0.1675 - val_acc: 0.9509 - val_loss: 0.1704


Epoch 13/30
81/81  **16s** 200ms/step - acc: 0.9551 - loss: 0.1487 - val_acc: 0.9418 - val_loss: 0.1754


Epoch 14/30
81/81  **17s** 203ms/step - acc: 0.9526 - loss: 0.1592 - val_acc: 0.9473 - val_loss: 0.1709


Epoch 15/30
81/81  **16s** 200ms/step - acc: 0.9546 - loss: 0.1557 - val_acc: 0.9491 - val_loss: 0.1625

Epoch 16/30
81/81  **17s** 204ms/step - acc: 0.9554 - loss: 0.1470 - val_acc: 0.9473 - val_loss: 0.1619

Epoch 17/30
81/81  **17s** 205ms/step - acc: 0.9544 - loss: 0.1392 - val_acc: 0.9491 - val_loss: 0.1666

Epoch 18/30
81/81  **16s** 200ms/step - acc: 0.9626 - loss: 0.1297 - val_acc: 0.9509 - val_loss: 0.1599

Epoch 19/30
81/81  **16s** 201ms/step - acc: 0.9612 - loss: 0.1341 - val_acc: 0.9509 - val_loss: 0.1577

Epoch 20/30
81/81  **17s** 205ms/step - acc: 0.9638 - loss: 0.1314 - val_acc: 0.9473 - val_loss: 0.1569

```

Epoch 21/30
81/81 ————— 16s 200ms/step - acc: 0.9532 - loss: 0.1344 - val_acc:
0.9382 - val_loss: 0.1667
Epoch 22/30
81/81 ————— 17s 203ms/step - acc: 0.9663 - loss: 0.1155 - val_acc:
0.9473 - val_loss: 0.1550
Epoch 23/30
81/81 ————— 17s 206ms/step - acc: 0.9607 - loss: 0.1188 - val_acc:
0.9455 - val_loss: 0.1556
Epoch 24/30
81/81 ————— 17s 210ms/step - acc: 0.9607 - loss: 0.1179 - val_acc:
0.9527 - val_loss: 0.1505
Epoch 25/30
81/81 ————— 17s 214ms/step - acc: 0.9618 - loss: 0.1184 - val_acc:
0.9473 - val_loss: 0.1581
Epoch 26/30
81/81 ————— 17s 211ms/step - acc: 0.9661 - loss: 0.1081 - val_acc:
0.9509 - val_loss: 0.1512
Epoch 27/30
81/81 ————— 17s 215ms/step - acc: 0.9622 - loss: 0.1162 - val_acc:
0.9473 - val_loss: 0.1507
Epoch 28/30
81/81 ————— 17s 203ms/step - acc: 0.9726 - loss: 0.0996 - val_acc:
0.9527 - val_loss: 0.1516
Epoch 29/30
81/81 ————— 17s 204ms/step - acc: 0.9652 - loss: 0.1010 - val_acc:
0.9509 - val_loss: 0.1482
Epoch 30/30
81/81 ————— 16s 200ms/step - acc: 0.9761 - loss: 0.0851 - val_acc:
0.9491 - val_loss: 0.1484

```

```

In [45]: # Assuming `model` is your trained Keras model
         model.save("CNN.h5") # Save in HDF5 format

```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

```

In [ ]: # Evaluate the model using the 'evaluate' method
         loss, acc = model.evaluate(test_batches, verbose=1)

         # Print results
         print("Loss: ", loss)
         print("Accuracy: ", acc)

```

```

18/18 ————— 3s 140ms/step - acc: 0.9433 - loss: 0.1720
Loss: 0.16484574973583221
Accuracy: 0.9490908980369568

```

```

In [16]: import numpy as np
         import matplotlib.pyplot as plt
         from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay

         # Evaluate the model using the 'evaluate' method
         loss, acc = model.evaluate(test_batches, verbose=1)

         # Print results
         print("Loss: ", loss)

```

```

print("Accuracy: ", acc)

# Generate predictions from the test data
y_pred = model.predict(test_batches)
y_pred_classes = np.argmax(y_pred, axis=1) # Convert predictions to class labels

# Get true labels from test_batches
y_true = test_batches.classes # Assuming test_batches is a generator with true labels

# Generate the confusion matrix
cm = confusion_matrix(y_true, y_pred_classes)

# Plot the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=test_batches.classes)
disp.plot(cmap=plt.cm.Blues)
plt.title("Confusion Matrix")
plt.show()


```

18/18 ————— 3s 140ms/step - acc: 0.9433 - loss: 0.1720

Loss: 0.16484574973583221

Accuracy: 0.9490908980369568

18/18 ————— 3s 141ms/step



****Diabet Retinopathy Detection Section****

```

In [ ]: import gradio as gr
import tensorflow as tf
from tensorflow.keras.preprocessing import image
import numpy as np

```



```

# Load the trained model
model = tf.keras.models.load_model(r'CNN.h5') # Replace with your model path

# Function to predict image
def predict_image(img):
    # Preprocess the image (resizing and scaling)
    img = img.resize((224, 224))
    img_array = np.array(img) / 255.0 # Rescale the image

    # Add batch dimension (as the model expects 4D input)
    img_array = np.expand_dims(img_array, axis=0)

    # Make prediction
    predictions = model.predict(img_array)

    # Get predicted class
    class_names = ['Diabetic Retinopathy', 'No Diabetic Retinopathy'] # Replace v
    predicted_class = class_names[np.argmax(predictions)]

    return predicted_class, predictions[0][np.argmax(predictions)]

# Build Gradio interface
iface = gr.Interface(fn=predict_image,
                    inputs=gr.Image(type="pil"),
                    outputs=[gr.Label(), gr.Textbox()],
                    live=True)

# Launch the Gradio app
iface.launch()

```

```

c:\Users\user\AppData\Local\Programs\Python\Python312\Lib\site-packages\tqdm\auto.p
y:21: TqdmWarning: IProgress not found. Please update jupyter and ipywidgets. See h
https://ipywidgets.readthedocs.io/en/stable/user_install.html

```

```

from .autonotebook import tqdm as notebook_tqdm
WARNING:absl:Compiled the loaded model, but the compiled metrics have yet to be bui
lt. `model.compile_metrics` will be empty until you train or evaluate the model.
* Running on local URL:  http://127.0.0.1:7860
* To create a public link, set `share=True` in `launch()`.

```



Clear

output0

Diabetic Retinopathy

output 1

Out[]:

1/1 ————— **2s** 2s/step

In []:

In []: